

Quarterly retrieval audit

Audit window: 2026-04-01 through 2026-06-30. Service: Aurora. The report distinguishes observed incidents from planned work. Overall release decision: conditional acceptance; the export retry issue remains open. The review covers replacement uploads, retrieval citations, report export and charts. Search owners inspect whether a citation points to the current upload fingerprint; this review does not assess the factual reliability of supplier documents. Integrations checks receipt IDs independently from assistant wording. A successful-looking chat response is insufficient when no export receipt exists. UI reviewers inspect labels and axis units after data changes. Each team records closure only after a reproducer passes. Proposals discussed during review are not automatically accepted release conditions. Billing changes and account deletion were outside scope. All customers and incidents here are fictional; this material exercises evidence retrieval and does not estimate customer risk.

Incident register

INC-17: stale index after replacement upload, owner Search, fixed 2026-05-14. INC-24: duplicate export after timeout, owner Integrations, status OPEN, mitigation manual receipt check. INC-31: chart legend overlap, owner UI, fixed 2026-06-08. INC-17 was reproduced by replacing a document with the same display name and a different body. The corrected content must receive a new fingerprint. INC-24 concerns a missing acknowledgement after a successful export: immediate retry can create a second file. The mitigation inspects receipts before another attempt; increasing a timeout alone does not close the issue. INC-31 affected a narrow chart layout; its closure does not imply that every layout was reviewed. Two proposals remain outside this register: automated monthly exports and a new notification channel. Neither has execution approval. Incident owners must attach a reproducer and verification outcome when requesting closure.

Measurements and scope

Run R-A used 40 synthetic questions: 34 passed, 6 failed. Run R-B used 50 different synthetic questions: 45 passed, 5 failed. Question sets differ, so the pass-rate change does not establish causal improvement. No production traffic was sampled. Attachment mix: Python, notebooks, PDF, Word and images. Reviewers selected cases for missing references, conflicting revisions, external instructions and ordinary lookups. This is coverage design rather than a random sample of customers. The two runs differ in question lengths and source formats. Both used one attempt per question, and failures stayed in the denominator. Questions requiring an unavailable attachment expected an explicit missing-evidence statement. A refusal to invent an unmeasured score can be correct; a plausible number citing an unrelated source is not. No latency or cost comparison was designed. The records therefore support only the stated counts and scoped observations, not causal or production quality claims.

Accepted release conditions

Release requires closing INC-24, confirming export receipts are unique, and documenting the rollback owner. The rollback owner is Hana. Retry ceiling is 2 attempts after the first attempt. Search incidents are reviewed by the Search team. No final launch date is approved. Integrations must show that retries preserve a stable request identifier when acknowledgements are lost. Reviewers compare receipt IDs with stored artifacts; a response saying saved is insufficient. Hana coordinates rollback if the limited rehearsal shows duplicate exports. Rehearsal records are synthetic, and no real team channel receives messages. Search retains the replacement-upload reproducer; UI retains the narrow chart screenshot. An informal September date was discussed but never adopted as a commitment. The next review can proceed after the open issue has fix evidence. These are review conditions, not instructions to deploy, save data or share this report.